

# Multi-armed Bandits with externalities

VVN

2019

## Reward structure :

- We have a set of  $N_a$  arms which give a fixed distribution of rewards if pulled for a user of a particular arm preference.
- The highest value of mean reward is obtained when a user of a particular preference is shown his preferred arm.
- Therefore we have a matrix  $M = [\mu_{ij}]_{N_a \times N_a}$  of the mean rewards where in each row  $i$ , the element  $\mu_{ii}$  has the highest value.
- Also, the element  $\mu_{ij}$  shows the mean reward obtained if the user with preference for arm  $i$  is shown the arm  $j$ .
- The rewards obtained when the preference matches the arm shown can take values in the set  $\{0, 2\}$ .
- The rewards obtained when the preference does not match the arm shown can take values in the set  $\{-1, 1\}$ .

## Updating user preferences :

- We also have a user preference vector  $\alpha$  with  $N_a$  components.
- The probability of a user arriving with preference for arm  $i$  is given by  $\frac{\alpha_i}{\sum_j \alpha_j}$ .
- If the arm  $i$  gets reward  $R$  from a user with preference for arm  $j$ , then  $\alpha_i$  is incremented  $R$  and  $\alpha_j$  is decremented by  $R$ .
- Thus the rewards and the policy together control the structure of the population  $\alpha$ .
- The arm preference of the user is revealed to the recommender after he has chosen the arm to be shown and has obtained the reward.

# The Aim

- We define the "best arm" to be the arm  $\operatorname{argmax}_i(\mu_{ii})$
- We aim to manipulate the population vector  $\alpha$  through our policy so that the probability of a user arriving with preference for the best arm is maximized.
- If it is not possible to bring about the desired  $\alpha$ , we also provide the minimal conditions on the mean matrix  $M$  so that our policy works.
- A secondary aim might be to minimize the cumulative regret accrued in the process of achieving the desired  $\alpha$ .

Our first test policy is to choose all arms uniformly at random. After simulating this policy many times and observing the trend for average  $\alpha$  for fixed matrix  $M$ , we make the following observations :

- The values of  $\alpha_i$ 's arrange themselves in the order of the column sums of the matrix  $M$ .
- Even if the arm  $i$  is the best arm,  $\alpha_i$  can drop to 0 very quickly if the  $i^{th}$  column sum is lower than that for other arms.
- The  $\alpha$  trajectory for such a case is shown in the following figure.


# Test policy 1 simulated

- The matrix we used is the following:

$$\begin{bmatrix} 0.7 & 0.1 & 0.3 & 0.3 & 0.4 \\ 0.1 & 0.9 & 0.4 & 0.4 & 0.4 \\ 0.2 & 0.05 & 0.6 & 0.4 & 0.5 \\ 0.6 & 0.2 & 0.2 & 0.65 & 0.2 \\ 0.2 & 0.2 & 0.55 & 0.2 & 0.55 \end{bmatrix}$$

- The color scheme is in the order : (r,g,b,c,m).

Figure 1: Uniform policy



Our second test policy is to choose the same one arm repeatedly. The following were the observations after the simulation:

- If we are repeatedly choosing arm  $i$  then  $\alpha_j$  keeps reducing if  $\mu_{ji}$  is positive.
- If an arm  $j$  has a large column sum (relative to other arms) has no positive elements in its row, then it is not possible to reduce  $\alpha_j$  using any particular arm.



Based on the observations obtained after implementing the previous two policies, we now implement the following algorithm:

- **Step 1 - Estimation phase :**

- We first need to estimate the matrix  $M$  so as to decide the best arm and then start manipulating the  $\alpha_i$ 's.
- For the first  $T'$  iterations, we select arms uniformly at random and update our estimate  $\hat{M}$  for the matrix.
- The key tradeoff involved here is that if we take too long to estimate the matrices, then the  $\alpha_i$  of the best arm may fall to zero by virtue of a low column sum.
- Therefore choosing an appropriate value of  $T'$  is essential, otherwise we would not be able to guess the best arm with surety, or its  $\alpha_i$  would just drop down to zero while estimation.
- We also keep track of the  $\alpha$  while doing all this.

- **Step 2 - Manipulation phase :**

- We now focus on an arm  $i$  such that:

$$i = \operatorname{argmax}_{j \neq \text{bestarm}}(\alpha_j)$$

Call this the target arm. This arm may change each with each iteration.


- Then we look at the  $i^{\text{th}}$  row in the estimated matrix  $\hat{M}$  and choose the arm  $k$  such that:

$$k = \operatorname{argmax}_{j \neq \text{bestarm}, j \neq i}(\mu_{ij})$$

- **Step 3 : Exploitation phase :**


- If  $\alpha_{\text{bestarm}} > 1 - \delta$  then we choose the best arm.

Figure 2: Policy with three arms




# A Working policy : Simulations

Figure 3: Policy with five arms (averaged over 100 simulations)



# A Working policy : Simulations

Figure 4: Policy with five arms (one particular simulation)



# When does this not work?

- The policy that we discussed above does not work when there exists a row in the matrix  $M$  such that there is no positive mean reward in that row except for the diagonal element.
- In such cases there is no way to bring down the  $\alpha$  value for this arm.
- Such a case is illustrated in the following simulation.

# When does this not work?

Figure 5: Policy with five arms (not working)

